

## COURSE OUTCOME-2

### PROGRAM NO-1

**Aim:** Program to implement k-NN classification using any standard dataset available in the public domain and find the accuracy of the algorithm.

```
import seaborn as sns
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

df = pd.read_csv('/content/sonar_csv.csv')
df.head()

   attribute_1 attribute_2 attribute_3 ... attribute_59
attribute_60 Class
0      0.0200      0.0371      0.0428 ...      0.0090
0.0032   Rock
1      0.0453      0.0523      0.0843 ...      0.0052
0.0044   Rock
2      0.0262      0.0582      0.1099 ...      0.0095
0.0078   Rock
3      0.0100      0.0171      0.0623 ...      0.0040
0.0117   Rock
4      0.0762      0.0666      0.0481 ...      0.0107
0.0094   Rock

[5 rows x 61 columns]

from sklearn.preprocessing import StandardScaler
#Preprocessing- standardscaler

sc=StandardScaler()

sc.fit(df.drop('Class',axis=1))

StandardScaler()

s = sc.transform(df.drop('Class',axis=1))

s

array([[ -0.39955135, -0.04064823, -0.02692565, ...,  0.06987027,
         0.17167808, -0.65894689],
       [  0.70353822,  0.42163039,  1.05561832, ..., -0.47240644,
        -0.44455424, -0.41985233],
       [ -0.12922901,  0.60106749,  1.72340448, ...,  1.30935987,
         0.25276128,  0.25758223],
       ...,
       [  1.00438083,  0.16007801, -0.67384349, ...,  0.90652575,
```

```

        -0.03913824, -0.67887143],
    [ 0.04953255, -0.09539176,  0.13480381, ..., -0.00759783,
      -0.70402047, -0.34015415],
    [-0.13794908, -0.06497869, -0.78861924, ..., -0.6738235 ,
      -0.29860448,  0.99479044]])

a = pd.DataFrame(s,columns = df.columns[:-1])

a

```

	attribute_1	attribute_2	...	attribute_59	attribute_60
0	-0.399551	-0.040648	...	0.171678	-0.658947
1	0.703538	0.421630	...	-0.444554	-0.419852
2	-0.129229	0.601067	...	0.252761	0.257582
3	-0.835555	-0.648910	...	-0.639154	1.034640
4	2.050790	0.856537	...	0.447361	0.576375
...	...	...	...	...	...
203	-0.456232	-0.116681	...	1.841992	1.831621
204	0.136733	-0.861801	...	-0.282388	0.038412
205	1.004381	0.160078	...	-0.039138	-0.678871
206	0.049533	-0.095392	...	-0.704020	-0.340154
207	-0.137949	-0.064979	...	-0.298604	0.994790

```

[208 rows x 60 columns]

```

training and testing

```
from sklearn.model_selection import train_test_split
```

```
x = a
```

```
y = df['Class']
```

```
y
```

```

0      Rock
1      Rock
2      Rock
3      Rock
4      Rock

```

```

...
203    Mine
204    Mine
205    Mine
206    Mine
207    Mine

```

```
Name: Class, Length: 208, dtype: object
```

```

x_train,x_test,y_train,y_test =
train_test_split(x,y,test_size=0.3,random_state= 42)

```

```
from sklearn.neighbors import KNeighborsClassifier
```

```

knn = KNeighborsClassifier(n_neighbors=1)
knn.fit(x_train,y_train)
KNeighborsClassifier(n_neighbors=1)
pred = knn.predict(x_test)
pred
array(['Mine', 'Mine', 'Rock', 'Rock', 'Mine', 'Rock', 'Mine', 'Mine',
       'Rock', 'Rock', 'Mine', 'Rock', 'Mine', 'Mine', 'Mine', 'Mine',
       'Mine', 'Mine', 'Rock', 'Rock', 'Mine', 'Mine', 'Mine', 'Mine',
       'Rock', 'Rock', 'Rock', 'Rock', 'Mine', 'Mine', 'Mine', 'Rock',
       'Rock', 'Mine', 'Mine', 'Mine', 'Rock', 'Mine', 'Mine', 'Mine',
       'Rock', 'Mine', 'Rock', 'Mine', 'Rock', 'Rock', 'Mine', 'Mine',
       'Rock', 'Rock', 'Rock', 'Mine', 'Rock', 'Rock', 'Rock', 'Rock',
       'Mine', 'Mine', 'Rock', 'Mine', 'Rock', 'Mine', 'Rock'],
      dtype=object)

```

```

from sklearn.metrics import classification_report,confusion_matrix
print(confusion_matrix(y_test,pred))
print(classification_report(y_test,pred))

```

```

[[32  3]
 [ 2 26]]

```

	precision	recall	f1-score	support
Mine	0.94	0.91	0.93	35
Rock	0.90	0.93	0.91	28
accuracy			0.92	63
macro avg	0.92	0.92	0.92	63
weighted avg	0.92	0.92	0.92	63

*#analyzing better k value through iterations*

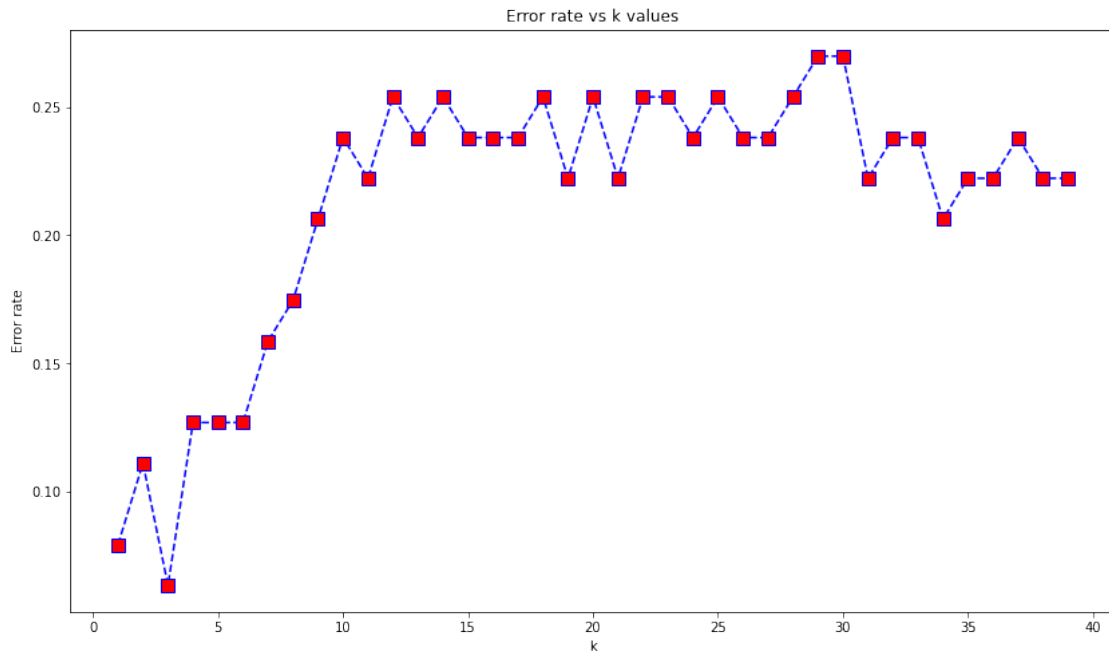
```

error_rate = []
for i in range(1,40):
    knn = KNeighborsClassifier(n_neighbors=i)
    knn.fit(x_train,y_train)
    pred_i=knn.predict(x_test)
    error_rate.append(np.mean(pred_i != y_test))

plt.figure(figsize=(14,8))
plt.plot(range(1,40),error_rate,color = 'blue',linestyle = 'dashed',
         marker= 's',markerfacecolor= 'red',markersize = 10)
plt.title('Error rate vs k values')
plt.xlabel('k')
plt.ylabel('Error rate')

```

```
Text(0, 0.5, 'Error rate')
```



```
knn = KNeighborsClassifier(n_neighbors=3)
knn.fit(x_train,y_train)
pred = knn.predict(x_test)
print(confusion_matrix(y_test,pred))
print(classification_report(y_test,pred))
```

```
[[33  2]
 [ 2 26]]
```

	precision	recall	f1-score	support
Mine	0.94	0.94	0.94	35
Rock	0.93	0.93	0.93	28
accuracy			0.94	63
macro avg	0.94	0.94	0.94	63
weighted avg	0.94	0.94	0.94	63

**Result:** The program is executed successfully and obtained the output.

## PROGRAM NO-2

**Aim:** Program to implement Naïve Bayes Algorithm using any standard dataset available in the public domain and find the accuracy of the algorithm

```
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
```

```
iris= sns.load_dataset('iris')
```


```
iris.head()
```

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	setosa
1	4.9	3.0	1.4	0.2	setosa
2	4.7	3.2	1.3	0.2	setosa
3	4.6	3.1	1.5	0.2	setosa
4	5.0	3.6	1.4	0.2	setosa

Pairgrid

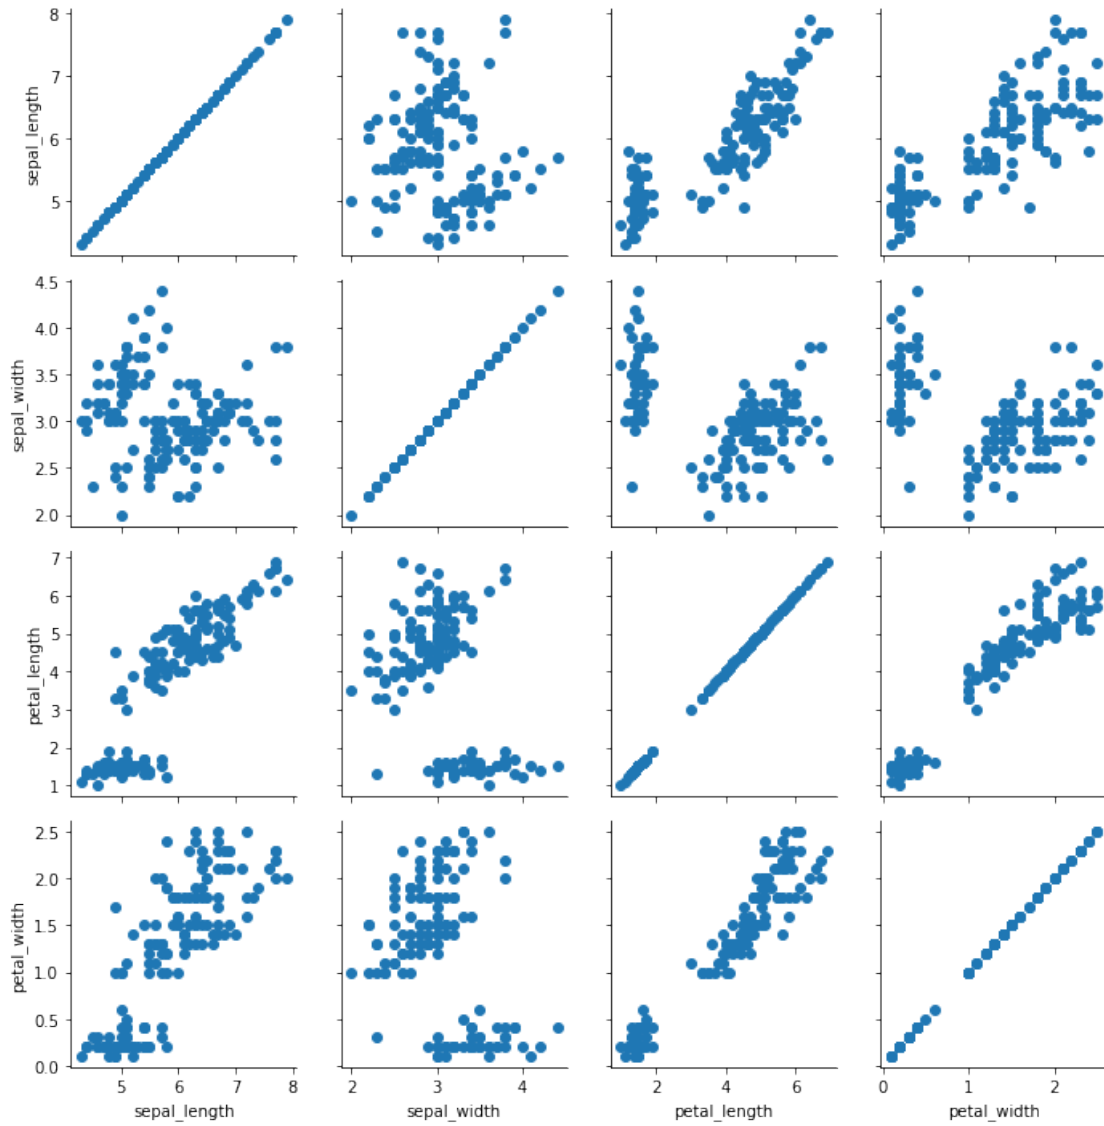
```
sns.PairGrid(iris)
```

```
<seaborn.axisgrid.PairGrid at 0x7f08b23f1810>
```



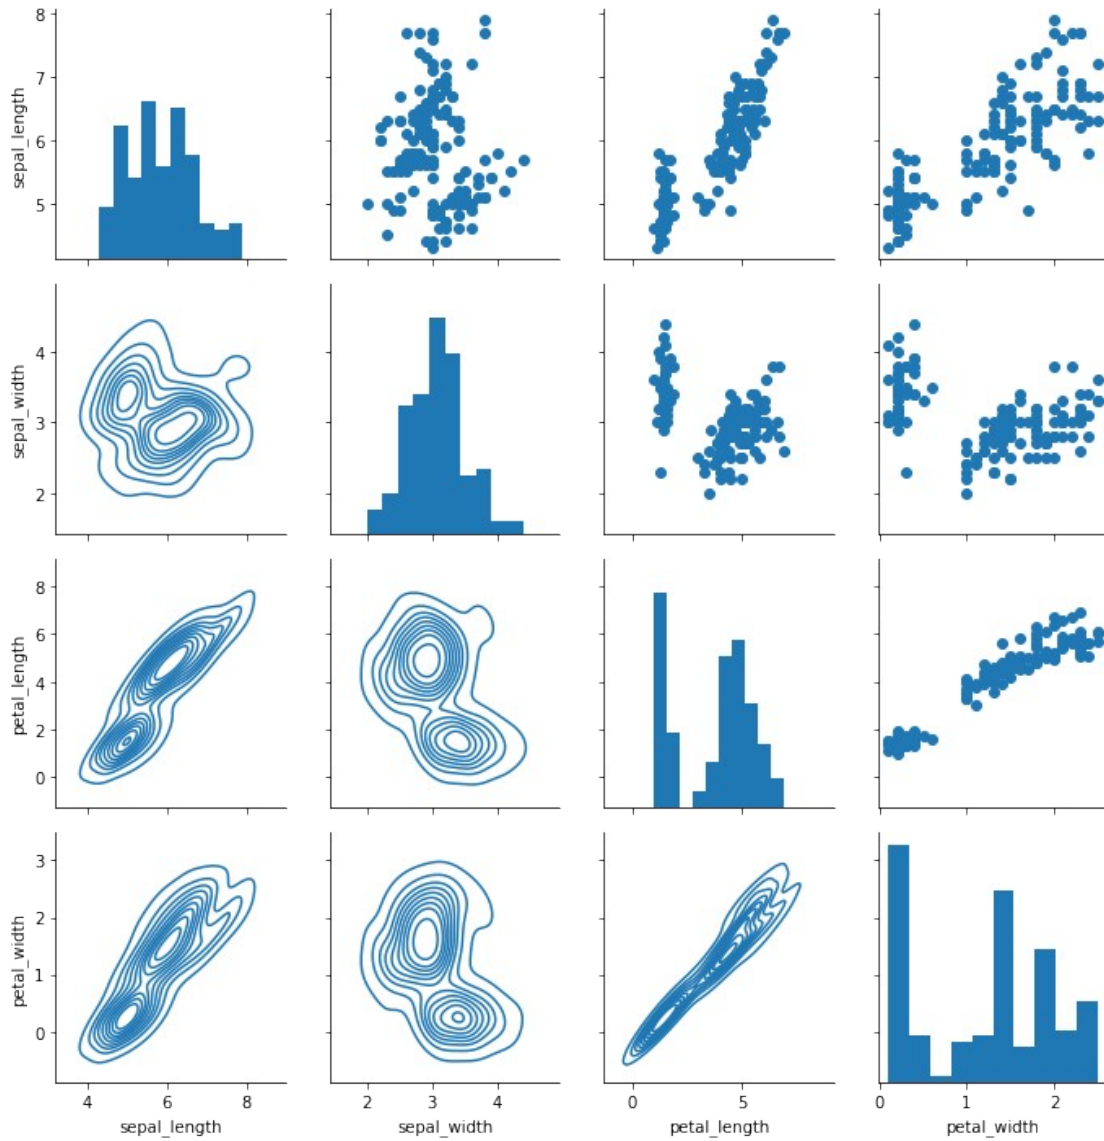
```
g = sns.PairGrid(iris)
g.map(plt.scatter) # mapping to the grid

<seaborn.axisgrid.PairGrid at 0x7f08a91661d0>
```



```
g = sns.PairGrid(iris)
g.map_diag(plt.hist)
g.map_upper(plt.scatter)
g.map_lower(sns.kdeplot)
```

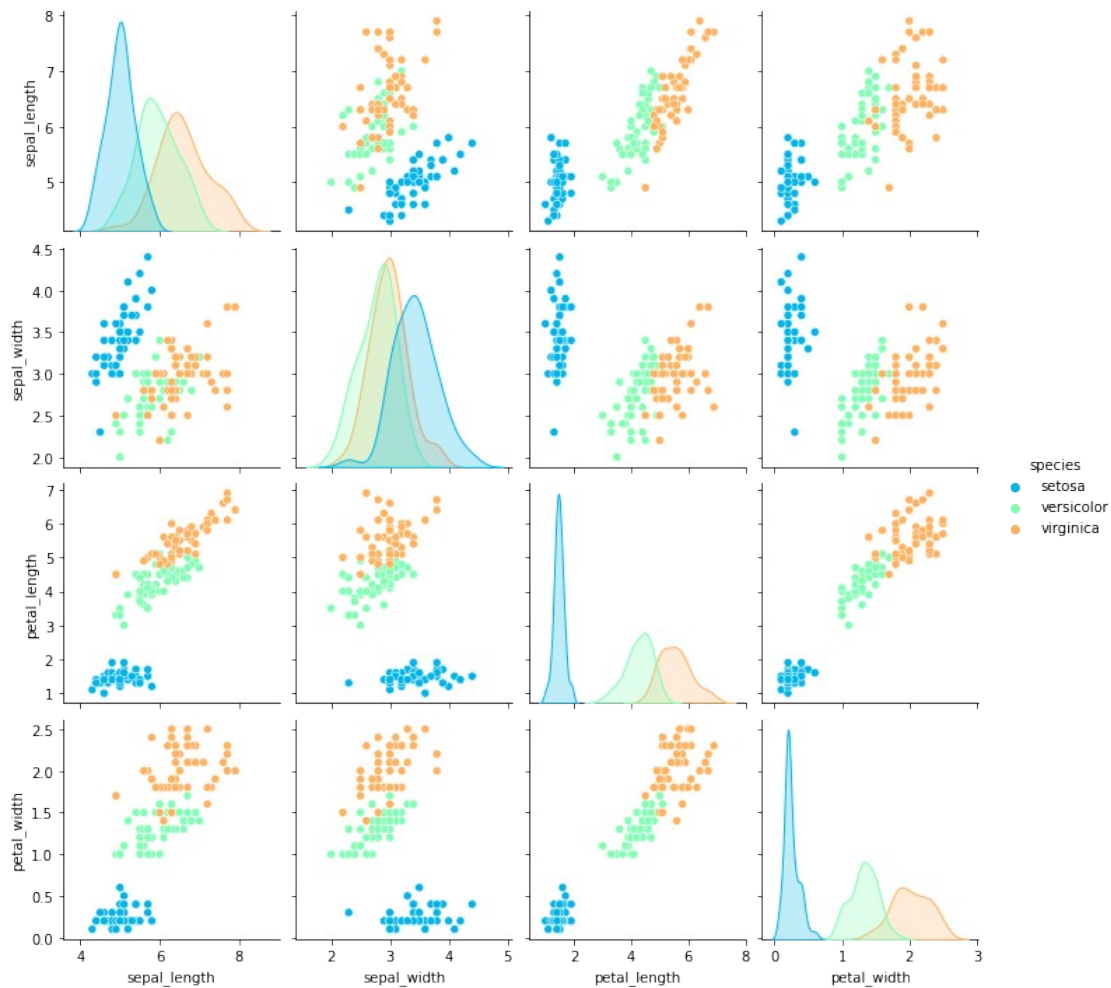
<seaborn.axisgrid.PairGrid at 0x7f08b1ee85d0>



```
sns.pairplot(iris,hue='species',palette ='rainbow')
```

```
<seaborn.axisgrid.PairGrid at 0x7f08a3035090>
```






facetgrid


```
tips = sns.load_dataset('tips')
tips.head()
```

	total_bill	tip	sex	smoker	day	time	size
0	16.99	1.01	Female	No	Sun	Dinner	2
1	10.34	1.66	Male	No	Sun	Dinner	3
2	21.01	3.50	Male	No	Sun	Dinner	3
3	23.68	3.31	Male	No	Sun	Dinner	2
4	24.59	3.61	Female	No	Sun	Dinner	4

```
g = sns.FacetGrid(tips, col= 'time', row='smoker')
```




```
g = sns.FacetGrid(tips, col= 'time',row='smoker')  
g =g.map(plt.hist,'total_bill')
```



Regression Plot


```
sns.lmplot(x='total_bill',y='tip',data=tips)
```

```
<seaborn.axisgrid.FacetGrid at 0x7f08a249dd10>
```




```
sns.lmplot(x='total_bill',y='tip',hue='sex',data=tips)
```

```
<seaborn.axisgrid.FacetGrid at 0x7f08a2427a50>
```



```
sns.lmplot(x='total_bill',y='size',hue='sex',data=tips)
```


```
<seaborn.axisgrid.FacetGrid at 0x7f08a23b0ad0>
```




Style **and** color

```
sns.countplot(x='sex',data = tips)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f08a234b990>
```




```
sns.set_style('dark')  
sns.countplot(x='sex',data = tips)  
<matplotlib.axes._subplots.AxesSubplot at 0x7f08a2286bd0>
```




```
sns.set_style('ticks')  
sns.countplot(x='sex',data = tips)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f08a226cad0>
```




```
sns.countplot(x='sex',data = tips)  
sns.despine()
```




```
sns.countplot(x='sex',data = tips)  
sns.despine(left = True)
```





```
plt.figure(figsize=(12,3))  
resns.countplot(x='sex',data = tips)  
<matplotlib.axes._subplots.AxesSubplot at 0x7f08a1fee450>
```



**Result:** The program is executed successfully and obtained the output.

### PROGRAM NO-3

**Aim:** Program to implement linear and multiple regression techniques using any standard dataset available in the public domain and evaluate its performance

*#Simple Linear Regression*

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
df=pd.read_csv('/content/sample_data/Salary_Data.csv')
```

```
df.head()
```

	YearsExperience	Salary
0	1.1	39343.0
1	1.3	46205.0
2	1.5	37731.0
3	2.0	43525.0
4	2.2	39891.0

```
x=df.iloc[:, :-1].values
```

```
y=df.iloc[:, :-1].values
```

*# split the dataset into training set and testing set*

```
from sklearn.model_selection import train_test_split
x_test,x_train,y_test,y_train=train_test_split(x,y,test_size=0.3,random_state=30)
```

*# fit the training dataset*

```
from sklearn.linear_model import LinearRegression
regsr=LinearRegression()
regsr.fit(x_train,y_train)
```

```
LinearRegression()
```


```
y_pred=regsr.predict(x_test)
```

*#visualize*

```
plt.scatter(x_train,y_train,color='red')
plt.plot(x_train,regsr.predict(y_train),color='green')
plt.title("salary VS experience")
plt.xlabel("experience")
plt.ylabel("salary")
plt.show()
```



```
# visualizing the test results
plt.scatter(x_test,y_test,color='red')
plt.plot(x_test,regsr.predict(y_test),color='green')
plt.title("salary VS experience")
plt.xlabel("experience")
plt.ylabel("salary")
plt.show()
```



*# Multiple Linear Regression*

```
ds=pd.read_csv('/content/sample_data/combined cycle powerplant.csv')
```

```
ds.head()
```

	AT	V	AP	RH	PE
0	8.34	40.77	1010.84	90.01	480.48
1	23.64	58.49	1011.40	74.20	445.75
2	29.74	56.90	1007.15	41.91	438.76
3	19.07	49.69	1007.22	76.79	453.09
4	11.80	40.66	1017.13	97.20	464.43

```
u=ds.drop('PE',axis=1).values
```

```
v=ds['PE'].values
```

```
from sklearn.model_selection import train_test_split
```

```
u_train,u_test,v_train,v_test =
```

```
train_test_split(u,v,test_size=0.3,random_state=45)
```

*#fitting*

```
from sklearn.linear_model import LinearRegression
```

```
rgr=LinearRegression()
```

```
rgr.fit(u_train,v_train)
```

```
LinearRegression()
```

```
v_pred=rgr.predict(u_test)
```

```
print(v_pred)
```

```
[451.12275809 472.67973273 434.23317529 ... 479.20120415 470.80190333
 437.26990979]
```

```
#take the values of first row in u and compare our predicted v values
with actual v value
```

```
rgr.predict([[23.64 ,58.49, 1011.40, 74.20]])
```

```
array([445.23162503])
```


```
from sklearn.metrics import r2_score
r2_score(v_test,v_pred)
```

```
0.9270483924843018
```

```
# visualize
```

```
plt.figure(figsize=(15,10))
plt.scatter(v_test,v_pred)
plt.title("actual VS predicted")
plt.xlabel("actual")
plt.ylabel("predicted")
```

```
Text(0, 0.5, 'predicted')
```



```
#print prdedicted values of our model
```

```
pred_ds=pd.DataFrame({'Actual value':v_test , 'Predicted
value':v_pred})
pred_ds
```

	Actual value	Predicted value
0	449.23	451.122758
1	474.70	472.679733
2	434.18	434.233175
3	436.70	442.479946
4	477.27	481.164400
...	...	...
2866	465.26	462.625918
2867	441.71	442.161640
2868	477.51	479.201204
2869	467.62	470.801903
2870	438.52	437.269910

[2871 rows x 2 columns]

**Result:** The program is executed successfully and obtained the output.