

COURSE OUTCOME-3

PROGRAM NO-1

Aim: Program to implement text classification using Support vector machine

```
import nltk
import pandas as pd
```

```
nltk.download_shell()
```

NLTK Downloader

```
-----
-----
      d) Download   l) List     u) Update   c) Config   h) Help    q) Quit
-----
-----
Downloader> l
```

Packages:

```
[ ] abc..... Australian Broadcasting Commission 2006
[ ] alpino..... Alpino Dutch Treebank
[ ] averaged_perceptron_tagger Averaged Perceptron Tagger
[ ] averaged_perceptron_tagger_ru Averaged Perceptron Tagger
(Russian)
[ ] basque_grammars..... Grammars for Basque
[ ] biocreative_ppi..... BioCreAtIvE (Critical Assessment of
Information
                               Extraction Systems in Biology)
[ ] bllip_wsj_no_aux.... BLLIP Parser: WSJ Model
[ ] book_grammars..... Grammars from NLTK Book
[ ] brown..... Brown Corpus
[ ] brown_tei..... Brown Corpus (TEI XML Version)
[ ] cess_cat..... CESS-CAT Treebank
[ ] cess_esp..... CESS-ESP Treebank
[ ] chat80..... Chat-80 Data Files
[ ] city_database..... City Database
[ ] cmudict..... The Carnegie Mellon Pronouncing Dictionary
(0.6)
[ ] comparative_sentences Comparative Sentence Dataset
[ ] comtrans..... ComTrans Corpus Sample
[ ] conll2000..... CONLL 2000 Chunking Corpus
[ ] conll2002..... CONLL 2002 Named Entity Recognition Corpus
Hit Enter to continue: q
```

```
-----
-----
      d) Download   l) List     u) Update   c) Config   h) Help    q) Quit
-----
-----
```

Downloader> q

```
messages = [line.rstrip() for line in
open('/content/SMSSpamCollection')]
print(len(messages))
```

5574

messages[0]

```
{"type": "string"}
```

```
for mess_no, message in enumerate(messages[:10]):
    print(mess_no, message)
    print('\n')
```

0 hamGo until jurong point, crazy.. Available only in bugis n great world la e buffet... Cine there got amore wat...

1 hamOk lar... Joking wif u oni...

2 spam Free entry in 2 a wkly comp to win FA Cup final tkts 21st May 2005. Text FA to 87121 to receive entry question(std txt rate)T&C's apply 08452810075over18's

3 hamU dun say so early hor... U c already then say...

4 hamNah I don't think he goes to usf, he lives around here though

5 spam FreeMsg Hey there darling it's been 3 week's now and no word back! I'd like some fun you up for it still? Tb ok! XxX std chgs to send, £1.50 to rcv

6 hamEven my brother is not like to speak with me. They treat me like aids patent.

7 hamAs per your request 'Melle Melle (Oru Minnaminunginte Nurungu Vettam)' has been set as your callertune for all Callers. Press *9 to copy your friends Callertune

8 spam WINNER!! As a valued network customer you have been

selected to receive a £900 prize reward! To claim call 09061701461.
Claim code KL341. Valid 12 hours only.

9 spam Had your mobile 11 months or more? U R entitled to Update
to the latest colour mobiles with camera for Free! Call The Mobile
Update Co FREE on 08002986030

```
messages[0]
```

```
{"type": "string"}
```

```
import pandas as pd
```

```
messages = pd.read_csv('/content/SMSSpamCollection', sep='\n',  
names=['label', 'message'])  
messages.head()
```

```
   label message  
0  ham  Go until jurong point, crazy.. Available only ...  
1  ham                Ok lar... Joking wif u oni...  
2  spam Free entry in 2 a wkly comp to win FA Cup fina...  
3  ham  U dun say so early hor... U c already then say...  
4  ham  Nah I don't think he goes to usf, he lives aro...
```

```
#classification tasks needs numerical features, So converting strips  
into vector format using  
#1. function to split words from a sentence into list  
#2. remove stopwords
```

```
import string
```

SAMPLE CODE FOR REMOVING PUNCTUATIONS AND STOPWORDS:

```
#removing punctuations
```

```
mess = "Sample message! Notice: it has punctuation."  
string.punctuation
```

```
{"type": "string"}
```

```
nopunc = [c for c in mess if c not in string.punctuation]  
nopunc
```

```
['S',  
'a',  
'm',  
'p',  
'l',  
'e',  
' ',
```

```
'm',  
'e',  
's',  
's',  
'a',  
'g',  
'e',  
'',  
'N',  
'o',  
't',  
'i',  
'c',  
'e',  
'',  
'i',  
't',  
'',  
'h',  
'a',  
's',  
'',  
'p',  
'u',  
'n',  
'c',  
't',  
'u',  
'a',  
't',  
'i',  
'o',  
'n']
```

```
nopunc = ''.join(nopunc)  
nopunc
```

```
{"type": "string"}
```

```
#removing stopwords
```

```
#for this, we need to download stopword's corpus from nltk.corpus
```

```
import stopwords
```

```
from nltk.corpus import stopwords
```

```
import nltk
```

```
nltk.download('stopwords')
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
```

```
[nltk_data] Package stopwords is already up-to-date!
```

```
True
```

```

nopunc.split()

['Sample', 'message', 'Notice', 'it', 'has', 'punctuation']

clean_mess = [word for word in nopunc.split() if word.lower() not in
stopwords.words('english')]

clean_mess

['Sample', 'message', 'Notice', 'punctuation']

#apply above function in our actual dataset
def text_process(mess):
    nopunc=[char for char in mess if char not in string.punctuation]
    nopunc="".join(nopunc)
    return[word for word in nopunc.split() if word.lower() not in
stopwords.words("english")]

messages.head()

  label      message
0  ham  Go until jurong point, crazy.. Available only ...
1  ham                Ok lar... Joking wif u oni...
2  spam  Free entry in 2 a wkly comp to win FA Cup fina...
3  ham  U dun say so early hor... U c already then say...
4  ham  Nah I don't think he goes to usf, he lives aro...

#tokenize
messages['message'].head(5).apply(text_process)

0    [Go, jurong, point, crazy, Available, bugis, n...
1                [Ok, lar, Joking, wif, u, oni]
2    [Free, entry, 2, wkly, comp, win, FA, Cup, fin...
3        [U, dun, say, early, hor, U, c, already, say]
4    [Nah, dont, think, goes, usf, lives, around, t...
Name: message, dtype: object

#converting tokens into vectors so that our machine learning models
get understand
from sklearn.feature_extraction.text import CountVectorizer

bow_transformer=CountVectorizer(analyzer=text_process).fit(messages['m
essage'])

print(len(bow_transformer.vocabulary_))

11425

mess4=messages['message'][6]
print(mess4)

Even my brother is not like to speak with me. They treat me like aids
patent.

```

```
bow4=bow_transformer.transform([mess4])

print(bow4)

bow_transformer.get_feature_names()[7800]

/usr/local/lib/python3.7/dist-packages/sklearn/utils/
deprecation.py:87: FutureWarning: Function get_feature_names is
deprecated; get_feature_names is deprecated in 1.0 and will be removed
in 1.2. Please use get_feature_names_out instead.
  warnings.warn(msg, category=FutureWarning)
```

```
{"type": "string"}
```

```
#apply this transformation for the whole message column in the dataset
```

```
messages_bow=bow_transformer.transform(messages['message'])
print('shape of the Sparse Matrix:', messages['message'])

shape of the Sparse Matrix: 0          Go until jurong point, crazy..
Available only ...
1          Ok lar... Joking wif u oni...
2          Free entry in 2 a wkly comp to win FA Cup fina...
3          U dun say so early hor... U c already then say...
4          Nah I don't think he goes to usf, he lives aro...
...
5567       This is the 2nd time we have tried 2 contact u...
5568          Will ü b going to esplanade fr home?
5569       Pity, * was in mood for that. So...any other s...
5570       The guy did some bitching but I acted like i'd...
5571          Rofl. Its true to its name
Name: message, Length: 5572, dtype: object
```

```
#check how many nonzero occurrences
```

```
messages_bow.nnz
```

```
50548
```

```
TERM FREQUENCY-INVERSE DOCUMENT FREQUENCY
```

```
#difference between TF and DF
```

```
#TF is frequency counter for a term t in document d.
```

```
#DF is the count of occurrences of term t in the documents set N
```

```
from sklearn.feature_extraction.text import TfidfTransformer
```

```
tfidf_transformer=TfidfTransformer().fit(messages_bow)
```

```
tfidf4=tfidf_transformer.transform(bow4)
print(tfidf4)
```

```
(0, 10629)    0.3352766696931058
(0, 9971)     0.3268691780062757
(0, 8761)     0.43700993321905807
```

```
(0, 7800)      0.41453906826037096
(0, 5193)      0.33843411088434017
(0, 4590)      0.43700993321905807
(0, 1802)      0.3352766696931058
```

#converting the whole bag of words into tfidf

```
messages_tfidf = tfidf_transformer.transform(messages_bow)
```

```
from sklearn.naive_bayes import MultinomialNB
```

```
spam_detect_model=MultinomialNB().fit(messages_tfidf,messages['label']
)
```

```
all_pred = spam_detect_model.predict(messages_tfidf)
```

```
all_pred
```

```
array(['ham', 'ham', 'spam', ..., 'ham', 'ham', 'ham'], dtype='<U4')
```

```
from sklearn.model_selection import train_test_split
```

```
msg_train,msg_test,label_train,label_test =
train_test_split(messages['message'],messages['label'])
```

```
spam_detect_model=MultinomialNB().fit(messages_tfidf,messages['label']
)
```

```
predict = spam_detect_model.predict(messages_tfidf)
```

```
from sklearn.metrics import classification_report
print(classification_report(messages['label'],predict))
```

	precision	recall	f1-score	support
ham	0.98	1.00	0.99	4825
spam	1.00	0.85	0.92	747
accuracy			0.98	5572
macro avg	0.99	0.92	0.95	5572
weighted avg	0.98	0.98	0.98	5572

TRAIN TEST SPLIT

```
from sklearn.model_selection import train_test_split
msg_train,msg_test,label_test,label_train = \
train_test_split(messages['message'],messages['label'],test_size=0.2)
```

CREATING A DATA PIPELINE

#pipeline

```
from sklearn.pipeline import Pipeline
pipeline = Pipeline([
```

```

        ('bow',CountVectorizer(analyzer=text_process)),
#strings to token integer
        ('tfidf',TfidfTransformer()),
#integer counts to weighted TF-IDF score
        ('classifier',MultinomialNB()),
#train on TF-IDF vectors w/Naive Bayes
    ])

pipeline.fit(msg_test,label_train)

Pipeline(steps=[('bow',
                  CountVectorizer(analyzer=<function text_process at
0x7f1c627f2c20>)),
                ('tfidf', TfidfTransformer()),
                ('classifier', MultinomialNB())])

predictions = pipeline.predict(msg_test)
print(classification_report(predictions,label_train))

```

	precision	recall	f1-score	support
ham	1.00	0.95	0.97	1019
spam	0.64	1.00	0.78	96
accuracy			0.95	1115
macro avg	0.82	0.97	0.88	1115
weighted avg	0.97	0.95	0.96	1115

SVM CLASSIFIER

```

from sklearn import model_selection,naive_bayes,svm
from sklearn.metrics import accuracy_score

#classifier - algorithm - SVM
#fit the training dataset on the classifier
pipeline1 = Pipeline([
    ('bow',CountVectorizer(analyzer=text_process)),
#strings to token int
    ('tfidf',TfidfTransformer()),
#integer counts to weighted TF-IDF score

    ('classifier',svm.SVC(C=1.0,kernel='linear',degree=3,gamma='auto'))
])

pipeline1.fit(msg_test,label_train)

Pipeline(steps=[('bow',
                  CountVectorizer(analyzer=<function text_process at
0x7f1c627f2c20>)),

```



```

        ('tfidf', TfidfTransformer()),
        ('classifier', SVC(gamma='auto', kernel='linear'))])

predictions1 = pipeline1.predict(msg_test)

print(classification_report(predictions1,label_train))

```

	precision	recall	f1-score	support
ham	1.00	1.00	1.00	967
spam	0.99	1.00	1.00	148
accuracy			1.00	1115
macro avg	1.00	1.00	1.00	1115
weighted avg	1.00	1.00	1.00	1115

Result: The Program is executed successsfully and obtained the output.

PROGRAM NO-2

Aim: Program to implement decision trees using any standard dataset available in the public domain and find the accuracy of the algorithm.

```
import pandas as pd
import numpy as np
from sklearn.datasets import load_iris

data.data.shape

(150, 4)

#load iris data
data = load_iris()

print('classes to predict: ',data.target_names)
print('features: ',data.feature_names)

classes to predict: ['setosa' 'versicolor' 'virginica']
features: ['sepal length (cm)', 'sepal width (cm)', 'petal length (cm)', 'petal width (cm)']

x = data.data
y = data.target
display (x.shape,y.shape)

(150, 4)

(150,)

from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier

xtrain,xtest,ytrain,ytest = train_test_split(x,y,random_state = 50,
test_size = 0.25)

#default criterion is Gini
classifier = DecisionTreeClassifier()
classifier.fit(xtrain,ytrain)

DecisionTreeClassifier()

y_pred = classifier.predict(xtest)

from sklearn.metrics import accuracy_score
print('Accuracy on train data using Gini:',accuracy_score(ytrain,classifier.predict(xtrain)))
print('Accuracy on test data using Gini:',accuracy_score(ytest,y_pred))

Accuracy on train data using Gini: 1.0
Accuracy on test data using Gini: 0.9473684210526315
```

```
#change criterion to entropy
classifier_entropy = DecisionTreeClassifier(criterion = 'entropy')
classifier_entropy.fit(xtrain,ytrain)
y_pred_entropy = classifier_entropy.predict(xtest)
print('Accuracy on train data using entropy',
accuracy_score(ytrain,classifier.predict(xtrain)))
print('Accuracy on test data using entropy',
accuracy_score(ytest,y_pred_entropy))
```

Accuracy on train data using entropy 1.0
Accuracy on test data using entropy 0.9473684210526315

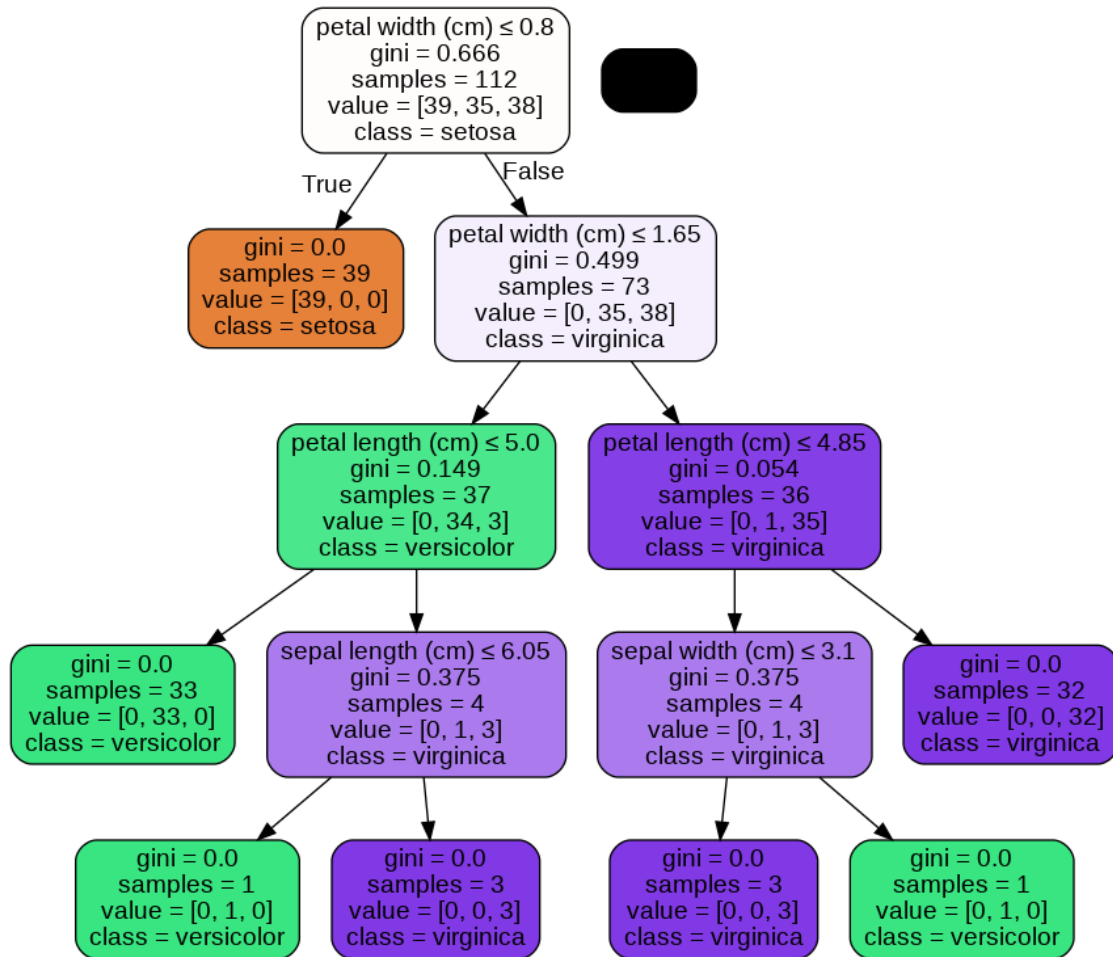
```
#change criterion to entropy with min_samples to 50. Default value is 2
classifier_entropy1 = DecisionTreeClassifier(criterion = 'entropy',
min_samples_split=50) #min_samples_split. min_samples_split represents
```

```
#the minimum number of samples required to split an internal node.
classifier_entropy1.fit(xtrain,ytrain)
y_pred_entropy1 = classifier_entropy1.predict(xtest)
print('Accuracy on train data using entropy',
accuracy_score(y_true=ytrain, y_pred=
classifier_entropy1.predict(xtrain)))
print('Accuracy on test data using entropy', accuracy_score(y_true =
ytest, y_pred = y_pred_entropy1))
```

Accuracy on train data using entropy 0.9642857142857143
Accuracy on test data using entropy 0.9473684210526315

```
#visualise the decision tree
from sklearn.tree import export_graphviz #for visualization
from six import StringIO #python 2,3 compatibility package, when the StringIO object is created
#it is initialized by passing a string to the constructor. If no string is passed the StringIO wil start empty.
from IPython.display import Image #Python is an interactive shell that is built with python
import pydotplus #Python interface to Graphviz's Dot language.
```

```
dot_data = StringIO()
#try using classifier,classifier_entropy and classifier_entropy1 as first parameter below.
export_graphviz(classifier, out_file = dot_data, filled = True,
rounded = True, special_characters = True, feature_names
=data.feature_names, class_names = data.target_names)
graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
Image(graph.create_png())
```



Result: The program is executed successfully and obtained the output.

PROGRAM NO-3

Aim: Program to implement k-means clustering technique using any standard dataset available in the public domain

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
df=pd.read_csv('/content/College_Data',index_col=0)
```

```
df.head()
```

	Private	Apps	...	Expend	Grad.Rate
Abilene Christian University	Yes	1660	...	7041	60
Adelphi University	Yes	2186	...	10527	56
Adrian College	Yes	1428	...	8735	54
Agnes Scott College	Yes	417	...	19016	59
Alaska Pacific University	Yes	193	...	10922	15

```
[5 rows x 18 columns]
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 777 entries, Abilene Christian University to York College of Pennsylvania
```

```
Data columns (total 18 columns):
```

#	Column	Non-Null Count	Dtype
0	Private	777 non-null	object
1	Apps	777 non-null	int64
2	Accept	777 non-null	int64
3	Enroll	777 non-null	int64
4	Top10perc	777 non-null	int64
5	Top25perc	777 non-null	int64
6	F.Undergrad	777 non-null	int64
7	P.Undergrad	777 non-null	int64
8	Outstate	777 non-null	int64
9	Room.Board	777 non-null	int64
10	Books	777 non-null	int64
11	Personal	777 non-null	int64
12	PhD	777 non-null	int64
13	Terminal	777 non-null	int64
14	S.F.Ratio	777 non-null	float64
15	perc.alumni	777 non-null	int64
16	Expend	777 non-null	int64
17	Grad.Rate	777 non-null	int64

```
dtypes: float64(1), int64(16), object(1)
memory usage: 131.5+ KB
```

```
df.describe()
```


	Apps	Accept	...	Expend	Grad.Rate
count	777.000000	777.000000	...	777.000000	777.000000
mean	3001.638353	2018.804376	...	9660.171171	65.46332
std	3870.201484	2451.113971	...	5221.768440	17.17771
min	81.000000	72.000000	...	3186.000000	10.00000
25%	776.000000	604.000000	...	6751.000000	53.00000
50%	1558.000000	1110.000000	...	8377.000000	65.00000
75%	3624.000000	2424.000000	...	10830.000000	78.00000
max	48094.000000	26330.000000	...	56233.000000	118.00000

```
[8 rows x 17 columns]
```

```
sns.set_style('darkgrid')
g=sns.FacetGrid(df,hue="Private",palette='coolwarm',size=6,aspect=2)
g=g.map(plt.hist,'Outstate',bins=20,alpha=0.7)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/axisgrid.py:337:
UserWarning: The `size` parameter has been renamed to `height`; please
update your code.
```

```
warnings.warn(msg, UserWarning)
```



```
from sklearn.cluster import KMeans
kmeans=KMeans(n_clusters=2)
kmeans.fit(df.drop('Private',axis=1))
KMeans(n_clusters=2)
kmeans.cluster_centers_
```

```
array([[1.03631389e+04, 6.55089815e+03, 2.56972222e+03,
4.14907407e+01,
       7.02037037e+01, 1.30619352e+04, 2.46486111e+03,
1.07191759e+04,
       4.64347222e+03, 5.95212963e+02, 1.71420370e+03,
8.63981481e+01,
       9.13333333e+01, 1.40277778e+01, 2.00740741e+01,
1.41705000e+04,
       6.75925926e+01],
 [1.81323468e+03, 1.28716592e+03, 4.91044843e+02,
2.53094170e+01,
       5.34708520e+01, 2.18854858e+03, 5.95458894e+02,
1.03957085e+04,
       4.31136472e+03, 5.41982063e+02, 1.28033632e+03,
7.04424514e+01,
       7.78251121e+01, 1.40997010e+01, 2.31748879e+01,
8.93204634e+03,
       6.51195815e+01]])
```

```
def converter(cluster):
    if cluster=='Yes':
        return 1
    else:
        return 0
```

```
df['cluster']=df['Private'].apply(converter)
```

```
df.head()
```

		Private	Apps	Accept	...	Expend
Grad.Rate	cluster					
Abilene Christian University	Yes	1660	1232	...	7041	
60	1					
Adelphi University	Yes	2186	1924	...	10527	
56	1					
Adrian College	Yes	1428	1097	...	8735	
54	1					
Agnes Scott College	Yes	417	349	...	19016	
59	1					
Alaska Pacific University	Yes	193	146	...	10922	
15	1					

```
[5 rows x 19 columns]
```

```
from sklearn.metrics import confusion_matrix,classification_report
print(confusion_matrix(df['cluster'],kmeans.labels_))
print(classification_report(df['cluster'],kmeans.labels_))
```

```
[[ 74 138]
 [ 34 531]]
```

```
precision    recall  f1-score   support
```

0	0.69	0.35	0.46	212
1	0.79	0.94	0.86	565
accuracy			0.78	777
macro avg	0.74	0.64	0.66	777
weighted avg	0.76	0.78	0.75	777

Result: The program is executed successfully and obtained the output.